

Lecture 4 : Lexical Analyzer using Flex

Compiler Design (CS 3007)


Sumanta Pyne

Assistant Professor
Computer Science and Engineering Department
National Institute of Technology Rourkela
pynes@nitrkl.ac.in

January 22, 2026


Lexical-Analyzer Generator Lex

- Lex tool = Lex language + Lex compiler
- More recent implementation is Flex
- Alternatives - JFlex uses Java, PyPI uses Python



- lex.l contains the regular expression
- lex.yy.c contains the lexical analyzer code in C language
- lex.yy.c is compiled to produce the lexical analyzer
- yylval - a global variable in lex.yy.c stores token attributes
- Token attribute - a numeric code, a pointer to symbol table or nothing
- yylval is shared between lexical analyzer and parser

Building a Compiler with Lex and Yacc



Structure of Lex Programs

- Format of a Lex program

```
declarations
%%
translation rules
%%
auxiliary functions
```

- declarations - variables, manifest constants and regular definitions
- translation rules - patterns with actions
- auxiliary functions - additional functions are used in the actions
- Format of each translation rule

Pattern { Action }

- Each pattern is a regular expression (may use regular definitions)
- Actions are fragment of code written in C

A Lex program to reconize following tokens

Token	Code	Value
begin	1	—
end	2	—
if	3	—
then	4	—
else	5	—
identifier	6	Pointer to symbol table
constant	7	Pointer to symbol table
<	8	1
<=	8	2
=	8	3
<>	8	4
>	8	5
>=	8	6

Tokens recognized

A Lex program - token_recognizer.l

```
%{
#include<stdio.h>
/* definitions of manifest constants */
enum yytokentype {LT = 1, LE, EQ, NE, GT, GE,
    BEGIN, END, IF, THEN, ELSE, ID, CONSTANT};
}%
/* regular definitions */
delim      [\t\n]
ws         {delim}+
letter     [A-Z,a-z]
digit      [0-9]
id         {letter}{(letter)|(digit)}*
constant   {digit}+
%%
{ws}       { /* no action and no return */}
begin      { printf("BEGIN\n"); return BEGIN;}
end        { printf("END\n"); return END;}
if         { printf("IF\n"); return IF;}
then       { printf("THEN\n"); return THEN;}
else       { printf("ELSE\n"); return ELSE;}
{id}       { printf("ID\n"); yylval = (int) installID(); return ID;}
{constant} { printf("CONSTANT\n"); yylval = (int) installConst();
    return CONSTANT;}
"<"       { printf("LT\n"); yylval = LT; return RELOP;}
"<="      { printf("LE\n"); yylval = LE; return RELOP;}
"="        { printf("LT\n"); yylval = EQ; return RELOP;}
"< >"     { printf("NE\n"); yylval = NE; return RELOP;}
">"       { printf("LT\n"); yylval = GT; return RELOP;}
">="      { printf("GE\n"); yylval = GE; return RELOP;}
%%
int installID(){ /* function to install the lexeme, whose
    first character is pointed to yytext,
    and whose length is yyleng, into the
    symbol table and return a pointer thereto*/
}
int installConst(){ /* similar to installID, but puts numerical
    constants into a separate table */
}
int main(){
    yylex(); /* scanner routine */
    return 0;
}
```

Running token_recognizer.l

```
$ flex token_recognizer.l
$ gcc lex.yy.c -lfl
$ ./a.out
if total > 50 then
IF
ID
GT
CONSTANT
THEN
begin 22 end
BEGIN
CONSTANT
END
^D
$
```

- lex.yy.c is linked with flex library, -lfl
- Each time the program needs a token, it calls yylex()
- yylex() reads an input, matches pattern and returns token
- yylex() called again for next token

Flex Library -lfl

- default main routine with `"while(yylex()!=0);"`
- `input()` - lexer calls `input()` to fetch each of the matched characters
- `unput()` - `unput(c)` returns character `c` to input stream
- `yyinput()` and `yyunput()` - `input()` and `unput()` in C++ scanner
- `yytext()` - matched token is stored in null-terminated string `yytext`
- `yytextlen()` - length of `yytext`
- `yyless(n)` - "push back" all but the first `n` characters of the token.
- `yylex()` and `YY_DECL` - `yylex` has no arguments, interacts through global variables. `YY_DECL` declares calling sequence and adds whatever arguments wanted for `yytext`
- `yywrap()` - tell lex to append the next token to current one
- `yyrestart()` - `yyrestart(f)` makes the scanner read from open stdio file `f`
- `yywrap()` - on reaching EOF lexer calls `yywrap()` to find next job, returns 0 for continuing scanning, returns 1 for EOF

Regular Expressions

- `[]` - any character within `[]` except `\n`
 - First character `^` - any character except the ones within `[]`
 - Dash in `[]` - character range
 - `[0 - 9]` means `[0123456789]`, `[a - z]` means any lowercase letter
 - `[a - z]{-}[jv]` - any character in $\{a, b, \dots, z\} - \{j, v\}$
- `$` - end of a line as the last character of a regular expression
- `{}` - min and max number of times the previous pattern can match
 - `A{1, 3}` matches one to three occurrences of the letter `A`
 - `0{5}` matches `00000`
- `\` - escape metacharacters as part of the usual C escape sequences
 - `\n` is a newline character, `\*` is a literal asterisk.
- `*` - zero or more copies of the preceding expression
- `+` - one or more copies of the preceding expression
- `?` - zero or one occurrence of the preceding regular expression
 - `-?[0 - 9]+` - signed number including an optional leading minus sign

Regular Expressions

- `|` - preceding or following regular expression
 - `faith|hope|charity` - any of the three virtues
 - `"..."` - anything within the quotation marks is treated literally
- `()` - groups a series of REs together into a new RE
 - `(01)` matches the character sequence 01
 - `a(bc|de)` matches *abc* or *ade*
- Few near misses in patterns
 - `[-+]?[0-9.]+` - matches too much, like 1.2.3.4
 - `[-+]?[0-9]+\.[0-9]+` - matches too little, misses .12 or 12.
 - `[-+]?[0-9]*\.[0-9]+` - doesn't match 12.
 - `[-+]?[0-9]+\.[0-9]*` - doesn't match .12
 - `[-+]?[0-9]*\.[0-9]*` - matches nothing, or a dot with no digits at all

How Flex Handles Ambiguous Patterns

- Ambiguous - multiple patterns that can match same input
- Match longest possible string every time scanner matches input
- Tie breaker - use pattern that appears first in the program
- For example - consider the following code snippet

```
"+" { return ADD; }
"=" { return ASSIGN; }
"+=" { return ASSIGNADD; }
"if" { return KEYWORDIF; }
"else" { return KEYWORDELSE; }
[a-zA-Z_][a-zA-Z0-9_]* { return IDENTIFIER; }
```
- += is matched as one token, since += is longer than +
- patterns for keywords precede patterns for identifier

File I/O in Flex Scanners

- A scanner reads from the stdio FILE called yyin
- To read a single file, set it before the first call to yylex
- Programs with simple I/O reads each input file from beginning to end
- In flex yyrestart(f) tells the scanner to read from stdio file f
- For multiple files
 - For each file open file
 - Use yyrestart() to make it the input to the scanner
 - Call yylex() to scan it

Word count program - one input file

- Word count, reading one file

```
%option noyywrap
%{
int chars = 0, words = 0, lines = 0;
}%
%%
[a-zA-Z]+      { words++; chars += strlen(yytext); }
\n             { chars++; lines++; }
.              { chars++; }
%%
main(int argc, char **argv){
    if(argc > 1) {
        if(!(yyin = fopen(argv[1], "r"))) {
            perror(argv[1]);
            return (1);
        }
    }
    yylex();
    printf("%8d%8d%8d\n", lines, words, chars);
}
```

Word count program - many input files

- Word count, reading many files

```
%option noyywrap % { int chars = 0, words = 0, lines = 0, totchars = 0, totwords = 0, totlines = 0; %} %%
[a - zA - Z]+ { words++; chars += strlen(yytext); }
\n { chars++; lines++; }
. { chars++; }
%%
main(int argc, char **argv){
    int i;
    if(argc < 2) { /* just read stdin */
        yylex(); printf("%8d%8d%8d\n", lines, words, chars);
        return 0;
    }
    for(i = 1; i < argc; i++) {
        FILE *f = fopen(argv[i], "r");
        if(!f) {
            perror(argv[i]);
            return (1);
        }
        yyrestart(f);
        yylex();
        fclose(f);
        printf("%8d%8d%8d %s\n", lines, words, chars, argv[i]);
        totchars += chars; chars = 0;
        totwords += words; words = 0;
        totlines += lines; lines = 0;
    }
    if(argc > 1) /* print total if more than one file */
        printf("%8d%8d%8d total\n", totlines, totwords, totchars);
    return 0;
}
```

The I/O Structure of a Flex Scanner

- Reads from an input source and optionally writes to an output sink
- By default, the input and output are stdin and stdout
- Lex read from yyin one character at a time
- Flex has three-level input allows user to handle any input structure
- Flex reads input using stdio from a file or standard input
- To handle input a flex uses a structure YY_BUFFER_STATE
- YY_BUFFER_STATE describes a single input source
- YY_BUFFER_STATE contains - a string buffer, variables and flags
- YY_BUFFER_STATE for scanning file or string already in memory
- Flex scanner output - by default copies unmatched input to yyout
 - . ECHO;
 - #define ECHO fwrite(yytext, yyleng, 1, yyout)

The I/O Structure of a Flex Scanner

- Default input behavior of a flex scanner:

```
YY_BUFFER_STATE bp;
```

```
extern FILE* yyin;
```

... whatever the program does before the first call to the scanner

if(!yyin) yyin = stdin; default input is stdin

```
bp = yy_create_buffer(yyin,YY_BUF_SIZE );
```

YY_BUF_SIZE defined by flex, typically 16K

*yy_switch_to_buffer(bp); tell it to use the buffer we just made yylex();
or yyparse() or whatever calls the scanner*

Thank you